Benchmark Radar v0.9.0

From daily collection to benchmark search and score history

Koutian Wu[1,2,*] · Junjie Zhou[3] · Jiayu Wang[4]
[1] Independent researcher
[2] Tacite AI
[3] Hangzhou Dianzi University
[4] Xi'an Jiaotong University
Corresponding author: Koutian Wu, k@tacite.ai
29 August 2026 | Software v0.9.0 | Data cutoff 2026-08-29 | Git 98c7de3
Reference DOI (frozen v0.9.0): 10.5281/zenodo.22167102

7,540	4,537	1,242	37
source observations across 37 snapshots	unique artifacts in the cumulative evidence graph	searchable entries across 4 search sources	public collection sources monitored

Benchmark Radar brings daily discoveries, catalog search, vendor reporting, and sourced model scores into one place. Each layer keeps its own count because each records a different kind of evidence.

Executive findings

- Search covers 1,242 entries. Three external catalogs supply 1,173 rows. The curated score archive adds 69 benchmark tracks.
- The daily corpus holds 7,540 observations for 4,537 artifacts. Primary or structured sources supplied 7,523 observations, and each record keeps its source URL and retrieval metadata.
- Thirty-six model reports cover 94 curated benchmarks. The same documents supply 285 scores across 69 benchmark tracks when the report includes a readable value and evaluation setup.
- Identity and task classification need more work. Forty-one artifacts link across multiple sources. KW-Bench has produced 4,129 tracks, with classification still empty in v0.9.0.

Abstract

Benchmark Radar collects benchmark work from 37 public sources and publishes it through a dashboard, RSS, JSON, and an offline client. This report checks the full path from collection to search. It covers source health, exact-identifier matching, the ranking rubric, the 1,242-entry search index, model-card adoption, score histories, and the reports already in the repository. The system has strong provenance and reproducible builds. Cross-source matching, protocol capture, semantic search, and KW-Bench classification remain active work.

1. Product tour

The README leads with the user task: find a benchmark in seconds, then inspect model-card adoption and score movement. The dashboard and RSS feed serve readers. JSON, the local CLI, and the local HTTP API support reproducible research.

Benchmark Radar's daily-intelligence logic was inspired by BuilderPulse, an AI-powered daily intelligence project for indie hackers and builders that answers 20 questions from 10+ sources every morning [8]. Benchmark Radar adapts that multi-source briefing pattern to benchmark discovery, normalized catalog search, and score history.

Search every benchmark

Search benchmarks, tasks, domains…

1,242 benchmarks · 4 sources

1.1 Search coverage

Search layer	Rows	Best field	Score data
LLM Stats	687	Benchmark names and descriptions	5,544 rows; evaluation setup is sparse
OpenCompass Hub	461	Paper, repository, release, and dataset links	Historical leaderboards stay source-labelled
Artificial Analysis	25	Current commercial evaluation catalog	7,050 third-party rows
Curated score tracks	69	Scores read from primary model reports	285 rows with instrument and protocol fields
Search box	**1,242**	**All four layers in one index**	**Source label shown on every result**

The UI adds 1,173 external rows and 69 curated score tracks. It keeps source records separate, so the same benchmark name can appear in more than one catalog. The label “4 sources” refers to these four search layers. The daily collector uses 37 public endpoints.

1.2 Dashboard and local tools

Surface	Reader question	Evaluation
Today	What appeared or changed?	Ranks new records and links the source evidence.
Search	Which records match this name, task, or domain?	Ranks exact names, phrases, and field-token matches.
Leaderboard	Which benchmarks do labs report?	Counts benchmark mentions across vendor reports.
Scores	Which printed values can be connected?	Matching instrument + protocol creates a series.
Trends / map	Which topics and sources recur?	Coverage signatures gate comparisons.
CLI + HTTP	Can an analyst reproduce search offline?	One QueryService and stable JSON contract.

2. Pipeline evaluation

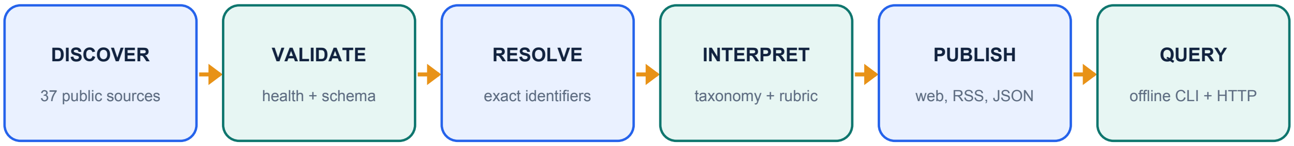

The same snapshots and code produce the same derived files.

2.1 Collection and health

Each run queries enabled connectors inside a 48-hour window, records counts and errors, drops future-dated rows, then requires the configured core sources to be healthy. arXiv, Hugging Face, and GitHub are required. On the cutoff run all three were healthy; Semantic Scholar returned HTTP 429 and Brave had no API key; OpenReview was healthy but empty.

The cutoff run fetched 826 rows, deduplicated them to 783, classified 273 as eligible, and recommended 97. Two collections ran that day. Their merged page contains 528 records, including 186 recommendations.

2.2 Normalization, identity, and retention

- Stable fields include source ID, URL, title, timestamps, authors or organizations when supplied, parser version, retrieval time, and a SHA-256 payload fingerprint. The publisher omits raw responses and credentials.
- The resolver merges observations that share a DOI, arXiv ID, OpenReview ID, GitHub repository, or Hugging Face artifact. A similar title leaves two records in place until a reviewer confirms the match.
- The corpus includes eligible records below the recommendation threshold. The recommended tag controls display priority.

2.3 Interpretation, publication, and query

Priority combines relevance (35%), evidence (20%), recency (20%), and adoption (25%) on a 0–100 scale. Scoring version 5 ships with the data. Generators replay validated snapshots into the cumulative graph, normalize external catalogs, classify KW-Bench tracks, and package a checksummed release. Installed clients switch to a new data version after checksum and schema validation. Search reads the active local version.

Area	Working today	Current gap
Reproducibility	Snapshots, schemas, parser versions, hashes, deterministic rebuilds.	Four early snapshots are simulated.
Reliability	Required-source gate; optional failures visible; atomic client activation.	Rate limits and secrets change realized coverage.
Identity	Exact-anchor merges and reviewed external groups.	Forty-one artifacts have more than one source.
Interfaces	CLI and HTTP share QueryService and JSON.	Local token matching misses some paraphrases.
Verification	Clean-worktree rebuild plus a passing full CI suite.	Source coverage still depends on public endpoints.

3. Counts and units

The dashboard publishes several counts because it tracks several units. Use the unit and cutoff beside the number.

Unit / product	Count	Definition	Use it for
Snapshot	37	33 observed, 4 simulated.	The dated history available for replay.
Source observation	7,540	Persisted source record.	Evidence volume in cumulative replay.
Unique artifact	4,537	Exact-ID-resolved paper, repo, dataset, release, or page.	Distinct artifacts observed.
External catalog row	1,173	One row from three catalog sources.	Broad, source-labelled catalog search.
Searchable entry	1,242	1,173 external + 69 curated score tracks.	Current web-search reach.
Adoption benchmark	94	Canonical identity in curated registry.	Model-card adoption and missing adoption.
Model-card document	36	Curated reports from 11 organizations.	Documents read for mentions.
Curated score	285	Numeric result from a cited document; 69 tracks.	Score histories grouped by evaluation setup.
Model registry	861	Models across curated and crawled layers; 19 in both.	Finding a model across both data layers.

3.1 Cumulative source composition

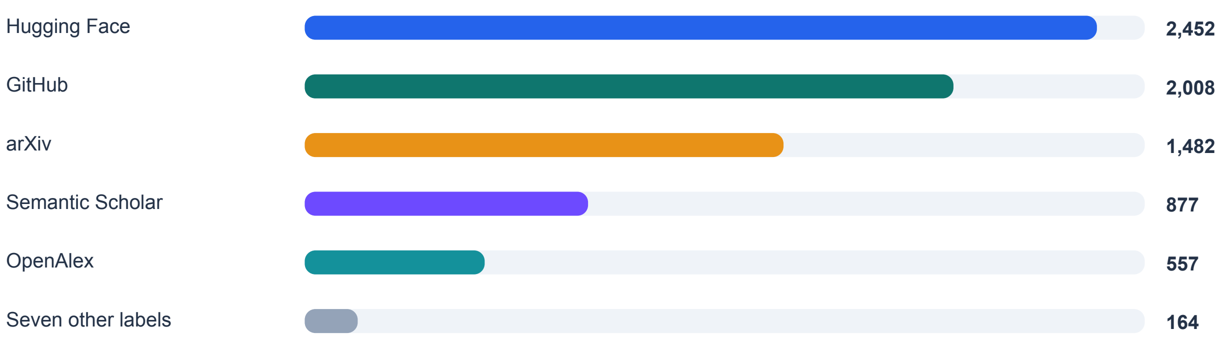

7,540 cumulative source observations through 29 August 2026

Five normalized source labels supply 7,376 of 7,540 observations (97.8%). The other endpoints add discovery routes and publish their own yield and health status.

3.2 Adoption and score findings

GPQA Diamond appears in 26 of 36 documents from 10 organizations. Humanity's Last Exam, SWE-bench Verified, Terminal-Bench, AIME, LiveCodeBench, MMLU-Pro, and BrowseComp form the rest of the top eight. These counts show which benchmarks vendors choose to report. The score archive finds eight bounded metrics with five points of headroom or less. Six benchmarks gained model-card mentions after their last readable score.

4. Source inventory and health

The daily collector reads 12 discovery connectors, 24 first-party feeds, and Hacker News. The search box draws from four catalog layers.

4.1 Direct discovery connectors (12)

Connector	Role	Cutoff run	Core?
arXiv	Primary papers via cs.AI/CL/CV/SE RSS	Healthy · 23	Yes
Hugging Face Hub	Datasets and Spaces	Healthy · 126	Yes
GitHub Search	Code and artifacts	Healthy · 300; at cap	Yes
GitHub Organizations	Reviewed organization repos	Healthy · 15	No
Hugging Face Papers	Community-surfaced papers	Healthy · 23	No
Kaggle Datasets	Public benchmark datasets	Healthy · 28	No
Zenodo	DOI-bearing artifacts	Healthy · 77	No
OpenReview	Conference submissions	Healthy · 0	No
Semantic Scholar	Structured scholarly discovery	Failed · HTTP 429	No
GitHub Releases	Curated first-party releases	Healthy · 3	No
OpenAlex	Scholarly discovery	Healthy · 228	No
Brave Search	Web and official domains	Unavailable · no key	No

4.2 First-party feeds (24)

Feeds 1–12	Feeds 13–24
Meituan Engineering OpenAI News Google AI Google DeepMind Google Research Apple Machine Learning Research AWS Machine Learning Hugging Face Blog Microsoft Research NVIDIA AI Blog Mistral AI Meta Research	Ai2 Together AI Sakana AI Qwen Ollama Stability AI Nomic AI Replicate NVIDIA Developer IBM Research Databricks LangChain

The feed collector yielded three relevant records. Some healthy feeds produced zero matches after relevance filtering. Domain-constrained searches cover organizations that lack a verified feed. The collector excludes third-party mirrors.

4.3 Public attention (1)

The Hacker News collector uses the public Algolia API. It found 12 records on the cutoff run and places them in a separate attention feed. The priority calculation uses evidence sources only.

5. Data quality and report history

Finding	Evidence	Interpretation
High provenance	7,523 / 7,540 primary or structured (99.77%).	Readers can open the upstream record for almost every observation.
Low corroboration	41 / 4,537 artifacts have >1 normalized source (0.90%).	Exact identity prevents bad joins but leaves plausible duplicates.
Uneven yield	Five source labels provide 97.8% of observations.	Caps or outages can change apparent topic mix.
Simulation	23–26 July are simulated snapshots.	Use them to test replay and UI history.
Classification gap	KW-Bench: 4,129 tracks, 0 classified.	The task-capability view has no classifications yet.
Protocol sparsity	12,594 aggregator scores stay outside curated progression.	Comparison needs shots, harness, tools, attempts, and date.
Lexical retrieval	Field tokens drive rank; no semantic reranker.	Reproducible, but paraphrases can be missed.

5.1 Existing reports

Document	Use today	Date limit
Landscape report, 31 Jul	A dated study of 791 sightings, 645 artifacts, and 78 agentic-evaluation artifacts.	Its totals stop on 31 July, before the external catalog and newer connectors.
Source probes, 27 Aug	Verified HF Papers, Kaggle, Spaces, and Zenodo endpoints.	The probes record one check on 27 August.
External catalog audit	OpenCompass identity fields and the score-heavy peer catalogs.	The audit keeps same-name records separate until identity review.
Daily report / briefing	Triage with citations and source health.	Each edition covers one collection window.

5.2 Reading the evidence

- You can trace a new or updated record to its source, inspect why the taxonomy matched it, and find the vendor document behind each curated benchmark mention.
- For score comparisons, match the instrument and protocol fields. For trend comparisons, use days with the same connector coverage and report limit. Market-size estimates need a separate benchmark-family census.

6. Findings from the current data

Benchmark Radar turns papers, repositories, datasets, and model reports into records you can search, filter, download, and query offline. Three results stand out in the current data.

6.1 Vendor attention has converged on a small reporting core

Eight benchmarks appear in documents from at least six organizations: GPQA Diamond, Humanity's Last Exam, SWE-bench Verified, Terminal-Bench, AIME, LiveCodeBench, MMLU-Pro, and BrowseComp. Teams comparing new model reports will encounter this group most often. The score archive shows where these familiar tests have little headroom left.

6.2 All eight raw near-ceiling readings come from one-date setups

All eight raw best readings leave five points of headroom or less, and every exact raw-best instrument+protocol setup appears on only one date. Four benchmarks have a different setup repeated across dates. Only HMMT's closest repeated setup is also within five points, and all four repeated setups are two-date, single-organization pairs. The remaining four benchmarks have no repeated setup. The evidence supports protocol-specific headroom claims only.

Benchmark	Raw gap	Same setup	Repeat gap	Claim decision	Note
AIME	0.8	1 date	20.2	**replace**	2 dates · 1 org
Arena-Hard	4.4	1 date	7.7	**replace**	2 dates · 1 org
DeepSearchQA	5	1 date	n/a	**qualify**	no repeated setup
HMMT	3.1	1 date	4.8	**qualify**	2 dates · 1 org
MATH-500	2	1 date	n/a	**qualify**	no repeated setup
MathVision	2.2	1 date	n/a	**qualify**	no repeated setup
SWE-bench Verified	4	1 date	19.4	**replace**	2 dates · 1 org
tau2-bench	0.7	1 date	n/a	**qualify**	no repeated setup

Threshold	Raw readings	Repeated setups
<=5	8 / 8	1 / 4 (4 unknown)
<=3	4 / 8	0 / 4 (4 unknown)
<=2	3 / 8	0 / 4 (4 unknown)

Repeated-setup sensitivity is 1 of 4 assessable benchmarks at five points; using 8 as the denominator would count four unknowns as negatives. No raw-best setup qualifies as repeat-controlled. The machine-readable audit beside the report source records every protocol stratum, score ID, exclusion, counterexample, and exact decision wording.

Per-benchmark claim wording

Benchmark	Exact report wording
AIME	AIME's 99.2 raw best leaves 0.8 points of headroom under the AIME 2026, temp 1.0, top_p 0.95, max gen 163840 tokens, GPT-5.5 (medium) judge setup, but that setup appears on one date. The closest repeated setup leaves 20.2 points, outside the five-point threshold.
Arena-Hard	Arena-Hard's 95.6 raw best leaves 4.4 points of headroom under the thinking mode setup, but that setup appears on one date. The closest repeated setup leaves 7.7 points, outside the five-point threshold.
DeepSearchQA	DeepSearchQA's 95 result leaves 5 points of headroom under the F1, reasoning=max setup. No setup spans two dates, so the archive supports only this protocol-specific reading.
HMMT	HMMT's raw best is an isolated 96.9 result. A separate think max, pass@1 setup spans 2 dates from 1 organization and leaves 4.8 points of headroom. This supports a protocol-specific paired comparison.

Benchmark	Exact report wording
MATH-500	MATH-500's 98 result leaves 2 points of headroom under the thinking mode setup. No setup spans two dates, so the archive supports only this protocol-specific reading.
MathVision	MathVision's 97.8 result leaves 2.2 points of headroom under the reasoning=max, with Python setup. No setup spans two dates, so the archive supports only this protocol-specific reading.
SWE-bench Verified	SWE-bench Verified's 96 raw best leaves 4 points of headroom under the adaptive thinking, max effort, averaged over 5 trials setup, but that setup appears on one date. The closest repeated setup leaves 19.4 points, outside the five-point threshold.
tau2-bench	tau2-bench's 99.3 result leaves 0.7 points of headroom under the Thinking (High), Telecom setup. No setup spans two dates, so the archive supports only this protocol-specific reading.

6.3 Broad search, deeper curation

The external catalog holds 1,173 rows, more than twelve times the 94-benchmark adoption registry. Use catalog search to find candidates. The curated registry adds the model reports, organizations, instruments, and protocols needed for comparison.

6.4 What to build next

Priority	Measurement work	Result
1	Show the count unit beside every UI total and export.	Readers can see which population each number counts.
2	Resolve high-value cross-source identities with anchors and review.	More records gain paper, code, and dataset links.
3	Complete KW-Bench extraction and reviewed coverage.	Turns 4,129 unclassified tracks into a task map.
4	Expand protocol capture from newer model cards.	Closes the mention-versus-score recency gap.
5	Add optional semantic retrieval while retaining lexical reasons.	Finds paraphrases while retaining visible match reasons.

6.5 Selected benchmark-family signal on agent weaknesses

Scope	Result
Issue #455 selected sample	8 demonstrated families; 6/8 state-control; 2/8 decision-execution; 4/4 agreement matches across 4 blinded sampled rows

Across 8 demonstrated benchmark families in the issue #455 selected sample, the family-deduplicated denominator shows a state-control-heavy pattern: 6/8 families fall in the coarse state-control grouping, versus 2/8 in decision-execution. This bounded selected sample is not a field-wide prevalence estimate.

Method and limits: Junkai Wang / @JunkaiWang-TheoPhy coded issue #455 (https://github.com/ktwu01/benchmark-radar/issues/455) from data/agent_weakness_evidence.yml at snapshot date 2026-09-01, evidence cutoff 2026-09-01, and commit input 98c8cf6fb5d1d69c66d438ea9f92242b2205c9ae. The demonstrated denominator excludes 1 design-implied row and 1 unmeasured row; each included family needed a primary-source anchor, a same-family counterexample, and family deduplication. Independent secondary coding matched on 4/4 completed blinded sampled rows out of 4 sampled rows; the 4/4 sample-local result among completed rows in the main packet does not establish broad reliability. SciCode remains the instrument counterexample rather than demonstrated prevalence evidence, because its audit shows that benchmark correction can reverse an apparent verification/completion failure signal. Primary-source benchmark evidence and the SciCode audit are cited in [9-17]. The sample is benchmark-family selected, not exhaustive, and limited by what current public benchmarks measure.

6.6 Worked real use case: prior-art check for a new evaluation

Jiayu Wang, a researcher working on agent evaluation, used Benchmark Radar to decide whether a proposed new evaluation would duplicate existing work. The check decides whether the design is still novel, and it used to be slow: comparing a candidate against the field required long manual searches, and completeness was hard to guarantee. The case ran during August 2026 with a concrete task: survey recent work on credit assignment in agentic training, keeping small Qwen-series baselines as a reproducibility constraint.

The author gave the task to a coding agent together with the public consumer prompt for Benchmark Radar. The agent installed the CLI and the benchmark-radar Skill, initialized the local corpus with benchmark-radar init, and queried candidate records with benchmark-radar search.

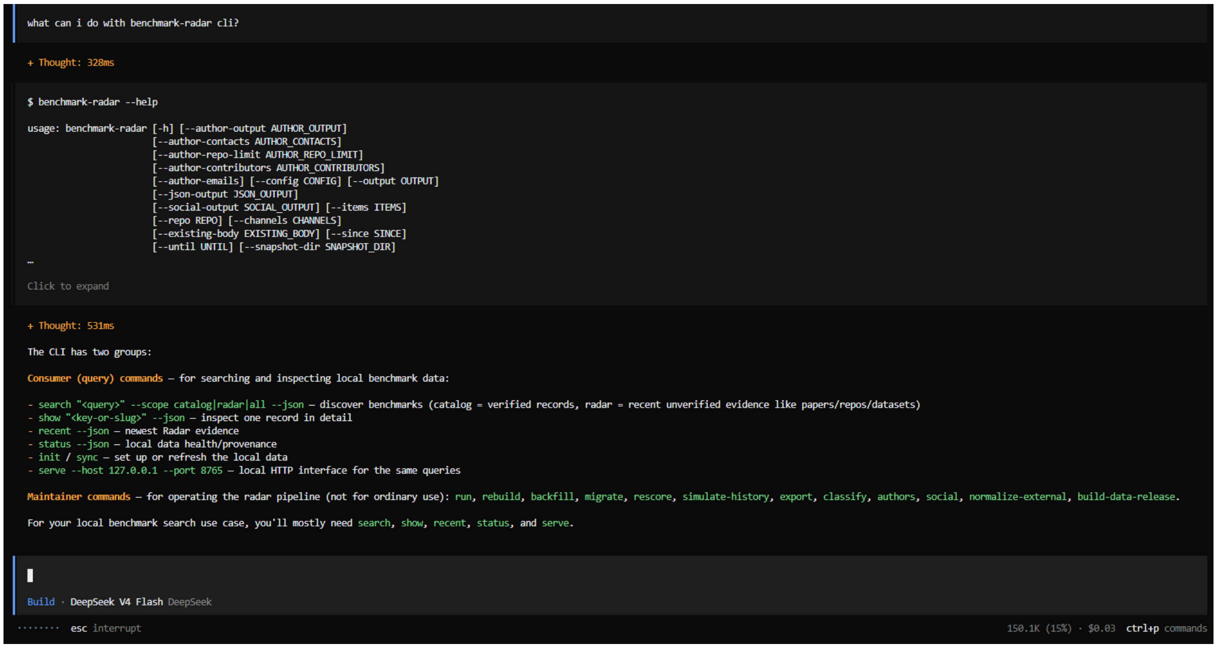

Figure 1. A coding agent follows the consumer setup prompt, installs the Benchmark Radar CLI and Skill, and runs local queries.

Radar links one artifact across papers, code, releases, and datasets. The agent could therefore see at a glance whether a candidate was announced as a paper with no released code, or shipped code without its dataset.

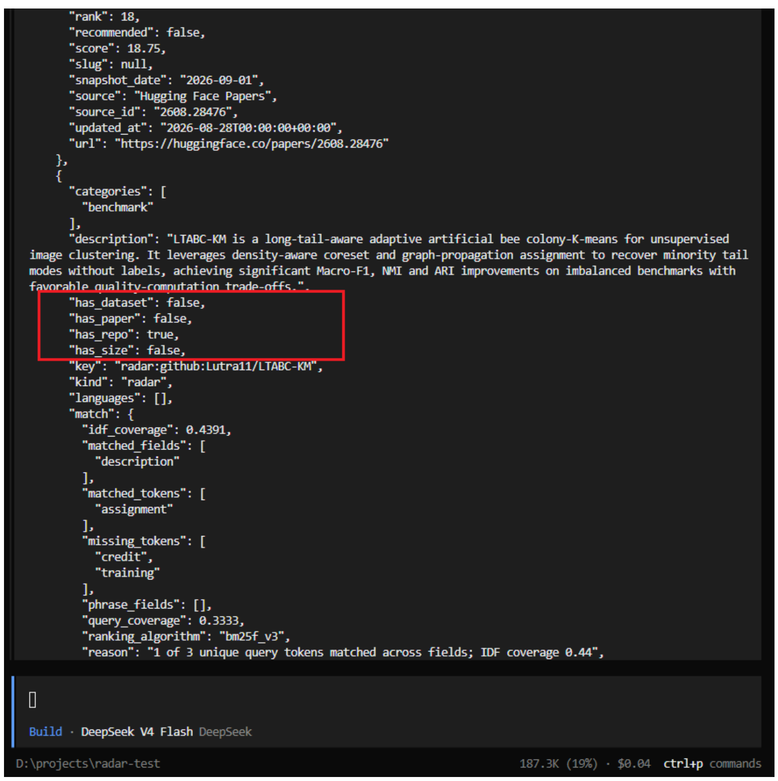

Figure 2. Radar consolidates the sources and status of one artifact.

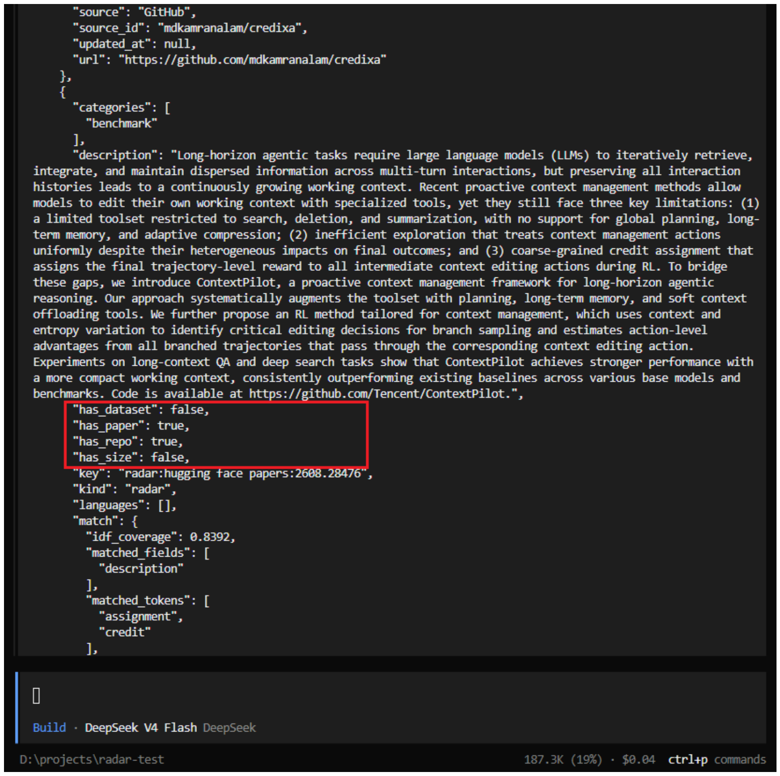

Figure 3. A companion record in which code is public but the dataset is not yet released.

Radar search is deterministic lexical matching, so the agent also ran its own web search and cross-checked the two candidate sets before accepting a record. This double pass keeps a differently worded version of the same idea from being missed.

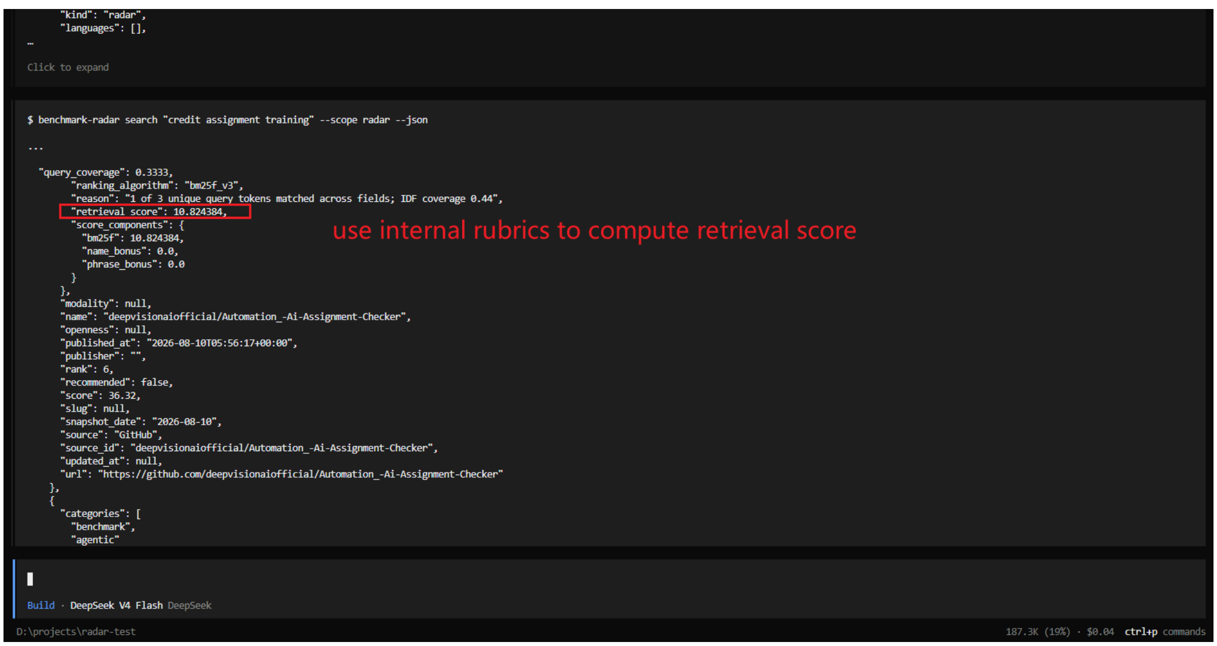

Figure 4. Cross-checking Radar candidates against the agent's own web search before accepting a record.

The session ended with a focused summary table of related work that the author judged complete enough to act on.

#	Work	Published	Qwen base model(s)	Credit-assignment focus	Agentic benchmarks used	Link
1	SRPO: Self-Reflective Policy Optimization for Long-Horizon Reasoning	2026-08-25	Qwen3-8B (also Qwen3-1.7B / 32B scaling)	Reflection-conditioned dense token-level signals convert sparse terminal reward into training signal	AIME'24, WebShop, ALFWorld, SWE-Bench-Lite	https://arxiv.org/abs/2608.23493 · https://github.com/Galleons2029/SRPO
2	ContextPilot: Teaching Agents for Proactive Context Management via Fine-grained RL	2026-08-28	Qwen3-8B, Qwen3-14B	Coarse-grained → action-level credit assignment over branched trajectories during agentic RL	InfBench (InfiniteBench), NovelQA, LongMemEval, BrowseComp+	https://arxiv.org/abs/2608.28476 · https://github.com/Tencent/ContextPilot
3	SkillGate: Training In-Policy Skill Selection in Long-Horizon Agents	2026-08-21	Qwen3.5-9B (SFT checkpoint, RL init)	"Selector credit starvation"; partitions token support into two credit channels (outcome vs. action-local advantage)	5 agentic benchmarks incl. SkillsBench, SWE (16-candidate skill slate)	https://arxiv.org/abs/2608.18852
4	CIPO: Contextual Information Policy Optimization for Search Agents	2026-08-06	Qwen2.5-3B-Instruct, Qwen2.5-7B-Instruct	Dense turn-level credit (EALR) to evidence-using reasoning actions, combined with outcome reward	HotpotQA, 2WikiMultiHopQA, MuSiQue, Bamboogle + 3 OOD	https://arxiv.org/abs/2608.06128
5	MoRSE: Task-Oriented Multi-Agent System with Mixture of Role-Subtask Experts	2026-08-10	Qwen3-4B-Instruct (one of three backbones; also Llama-3.1-8B, Gemma-4-31B)	Hierarchical GRPO with two-layer credit assignment (isolates expert vs. routing quality)	Code-generation benchmarks (multi-agent), held-out task domains	https://arxiv.org/abs/2608.09251

Figure 5. Summary table of recent work on credit assignment in agentic training assembled during the session.

The savings are easiest to measure against the author's earlier benchmark, AARRI-Bench, whose equivalent prior-art comparison consumed effort second only to producing the benchmark data itself, for a table of just 12 rows and 7 columns.

Table 1: **Comparison of Relevant AI research benchmarks.** AARRI-Bench simultaneously supports end-to-end research evaluation, fine-grained research process assessment, researcher quality and multi-harness evaluation.

Bench Name	End-to-End Tasks	Fine-Grained Eval	Researcher-Quality Eval	Data Generation	Multi-Harness Eval	#Tasks
MLE-Bench (Chan et al., 2025)	✗	✓	✗	Transfer&Compose	✓	75
MLGym-Bench (Nathani et al., 2025)	✗	✓	✗	Automatic	✗	13
EXP-Bench (Kon et al., 2025)	✓	✗	✗	Automatic	✓	461
ResearchCodeBench (Hua et al., 2026)	✗	✓	✗	Transfer&Compose	✗	212
MLR-Bench (Chen et al., 2026)	✓	✗	✗	Automatic	✓	201
PaperBench (Starace et al., 2025)	✗	✓	✗	Transfer&Compose	✗	8316
AstaBench (Bragg et al., 2025)	✓	✓	✗	Transfer&Compose	✓	2400+
InnovatorBench (Wu et al., 2025)	✓	✗	✗	Transfer&Compose	✗	20
AIRS-Bench (Lupidi et al., 2026)	✗	✓	✗	Automatic	✓	20
COMPOSITE-Stem (Waters et al., 2026)	✓	✓	✗	Manual	✗	70
ScienceBoard (Sun et al., 2025)	✗	✓	✗	Manual	✗	169
AARRI-Bench (Ours)	✓	✓	✓	**Manual**	✓	**82**

Figure 6. The manually built 12-by-7 prior-art table for AARRI-Bench, the workflow that Radar now shortens.

Limits: Radar search is deterministic lexical matching rather than semantic retrieval, so a differently worded query can change the candidate set. The session used the Radar CLI together with the agent's general web search, so it does not isolate Radar alone. Repository and dataset availability is a snapshot, not a permanent label. The case documents one contributor's workflow; it is not a measured user study. Full evidence, including the summary table and session screenshots, is public in issue #492.

Contributor. Jiayu Wang, Xi'an Jiaotong University. Case and evidence: github.com/ktwu01/benchmark-radar/issues/492

Use it Cite Benchmark Radar for source-linked benchmark discovery, catalog search, and vendor-reporting evidence. Use the source and protocol fields when you compare records or scores.

7. Reproducibility, access, and citation

This report evaluates Benchmark Radar v0.9.0 at Git commit 98c7de3 and data cutoff 2026-08-29. The clean worktree ran the CI sequence: lint and formatting checks, external normalization, KW-Bench classification, checksummed data-release construction, and the full test suite. The current full CI suite passed.

7.1 Methods and limitations

The 6.2 audit uses the curated score archive and model-card registry. Raw headroom is the distance from one reported value to the metric bound. Repeat-controlled evidence requires the same instrument and protocol across at least two dates. Two dates support only a paired comparison. Because each repeated setup comes from one organization, the audit cannot estimate a field-wide trend. Sensitivity tables keep unknowns out of the eligible denominator.

7.2 Contributor credit

Junjie Zhou prepared the protocol audit, machine-readable table, and report revision for issue #457.

7.3 Issue link

Issue #457 Near-ceiling metrics under protocol controls.

Artifact	Canonical or permanent location
Technical report	https://doi.org/10.5281/zenodo.22167102
Source code	https://github.com/ktwu01/benchmark-radar
Dashboard	https://benchmark-radar.org/
Cumulative JSON	https://benchmark-radar.org/data/radar.json
Benchmark catalog	https://benchmark-radar.org/data/benchmark-index.json
RSS	https://benchmark-radar.org/feed.xml
Citation metadata	https://github.com/ktwu01/benchmark-radar/blob/main/CITATION.cff

Published v0.9.0 citation

Wu, K. (2026). Benchmark Radar v0.9.0: Technical Report. Zenodo. https://doi.org/10.5281/zenodo.22167102

Data statement

The report's core counts are a frozen v0.9.0 audit at commit 98c7de3 with cutoff 2026-08-29, sourced from the versioned release files that back that snapshot. The current issue #455 study is reported separately as a 2026-09-01 selected-sample analysis and does not recompute or replace the frozen v0.9.0 core counts.

References

[1] K. Wu. Benchmark Radar, version 0.9.0. GitHub, 2026. https://github.com/ktwu01/benchmark-radar

[2] K. Wu. AI Benchmark Landscape Report. 2026. https://github.com/ktwu01/benchmark-radar/blob/main/docs/reports/ai-benchmark-landscape-report.md

[3] K. Wu. Benchmark Radar cumulative corpus schema. 2026. https://github.com/ktwu01/benchmark-radar/blob/main/docs/cumulative-corpus.schema.json

[4] K. Wu. Source probe evidence. 2026. https://github.com/ktwu01/benchmark-radar/blob/main/docs/source-probe-evidence.md

[5] D. S. Katz et al. Recognizing the value of software: a software citation guide. F1000Research 9:1257, 2021. https://doi.org/10.12688/f1000research.26932.2

[6] A. Smith, D. S. Katz, and K. E. Niemeyer. Software citation principles. PeerJ Computer Science 2:e86, 2016. https://doi.org/10.7717/peerj-cs.86

[7] Citation File Format developers. Citation File Format 1.2.0. https://citation-file-format.github.io/

[8] L. Xiaopai. BuilderPulse: AI-powered daily intelligence for indie hackers and builders. GitHub, 2026. https://github.com/BuilderPulse/BuilderPulse

[9] Primary-source evidence for OSWorld 2.0. https://arxiv.org/html/2606.29537v1. Evidence anchor: HTML abstract paragraph id abstract1.1, final two sentences on 20.6% binary completion, 54.8% partial score, and hidden-state recovery; Figure 8 caption at figure id S3.F8

[10] Primary-source evidence for WebArena. https://arxiv.org/html/2307.13854v4. Evidence anchor: HTML paragraph id S5.p1.1 on 11.70% success; HTML paragraph id S5.SS1.SSS0.Px1.p1.1 on 54.9%, 14.41%, and 44.44%; Table 2 caption at S5.T2

[11] Primary-source evidence for Mind2Web. https://arxiv.org/html/2306.06070v3. Evidence anchor: HTML paragraph id S4.SS3.p3.1 on the cross-task versus cross-website/domain gap and unseen-environment challenge; Figure 6 reference in the same paragraph

[12] Primary-source evidence for GAIA2. https://arxiv.org/html/2602.11964v1. Evidence anchor: HTML abstract paragraph id abstract1.1 on 42% pass@1 and time-sensitive failures; HTML paragraph id S5.SS2.p1.1 and Figure 8 caption id S5.F8 on default-versus-instant Time behavior

[13] Primary-source evidence for SWE-bench Verified. https://arxiv.org/abs/2509.13941. Evidence anchor: arXiv abs abstract paragraph on 150 failed instances, majority agentic failures from flawed reasoning and cognitive deadlocks, and 22.2% recovery

[14] Primary-source evidence for SWE-bench Science. https://arxiv.org/abs/2608.19799. Evidence anchor: arXiv abs abstract paragraph on pass@1 below 50%, the four recurring failure mechanisms, and the paired scientific-guidance ablation

[15] Primary-source evidence for ResearchClawBench. https://arxiv.org/html/2606.07591v5. Evidence anchor: HTML abstract paragraph on 21.5, 20.7, and protocol/evidence/scientific-core mismatch; HTML paragraph ids S4.SS2.p1.1 and S4.SS5.p1.1 under subsection 4.5 Error Analysis

[16] Primary-source evidence for PRBench. https://arxiv.org/html/2603.27646v1. Evidence anchor: HTML paragraph ids S4.SS2.p2.1 and S4.SS2.p4.1 on 34% overall score and 0% End-to-End Callback Rate; Table 2 caption at S4.T2

[17] Primary-source evidence for SciCode. https://arxiv.org/abs/2608.04975. Evidence anchor: arXiv abs abstract paragraph on 263 defects, 192 score-suppressing defects across 91% of main problems, and recovery to 84-98% / 69-92%